

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
get_ipython().run_line_magic('matplotlib', 'inline')
```

1. Loading Dataset Into Pandas Dataframe

```
In [3]: df = pd.read_csv(r'C:\Users\falco\OneDrive\Desktop\PROJECTS\PYTHON PROJECTS\MY PROJECT
```

2. Display Top 5 Rows()

```
In [4]: df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Nu
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	
3	4	15701354	Boni	699	France	Female	39	1	0.00	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	

3. Last 5 Rows

```
In [5]: df.tail()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	

4. Shape Of Dataset

```
In [8]: df.shape
```

```
print('Number of Rows: ', df.shape[0])
print('Number of Columns: ', df.shape[1])
```

Number of Rows: 10000
 Number of Columns: 14

In []: # 5. Overall Info About The Dataset

In [9]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   RowNumber         10000 non-null   int64  
 1   CustomerId        10000 non-null   int64  
 2   Surname           10000 non-null   object  
 3   CreditScore       10000 non-null   int64  
 4   Geography          10000 non-null   object  
 5   Gender             10000 non-null   object  
 6   Age                10000 non-null   int64  
 7   Tenure             10000 non-null   int64  
 8   Balance            10000 non-null   float64 
 9   NumOfProducts      10000 non-null   int64  
 10  HasCrCard          10000 non-null   int64  
 11  IsActiveMember     10000 non-null   int64  
 12  EstimatedSalary    10000 non-null   float64 
 13  Exited             10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In []: # 6. Null Values

In [10]: df.isnull().sum()

```
Out[10]: RowNumber      0
CustomerId      0
Surname        0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

7. Overall Statistics About The Dataset

In [11]: df.describe()

Out[11]:	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfP
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4

In [12]: `df.describe(include = 'all')`

Out[12]:	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
count	10000.00000	1.000000e+04	10000	10000.000000	10000	10000	10000.000000
unique	NaN	NaN	2932	NaN	3	2	NaN
top	NaN	NaN	Smith	NaN	France	Male	NaN
freq	NaN	NaN	32	NaN	5014	5457	NaN
mean	5000.50000	1.569094e+07	NaN	650.528800	NaN	NaN	38.921800
std	2886.89568	7.193619e+04	NaN	96.653299	NaN	NaN	10.487806
min	1.00000	1.556570e+07	NaN	350.000000	NaN	NaN	18.000000
25%	2500.75000	1.562853e+07	NaN	584.000000	NaN	NaN	32.000000
50%	5000.50000	1.569074e+07	NaN	652.000000	NaN	NaN	37.000000
75%	7500.25000	1.575323e+07	NaN	718.000000	NaN	NaN	44.000000
max	10000.00000	1.581569e+07	NaN	850.000000	NaN	NaN	92.000000

8. Dropping Irrelevant Columns

In [13]: `df.columns`

Out[13]: `Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'], dtype='object')`

In [15]: `df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis = 1)`

In [16]: `df.head()`

Out[16]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMe
0	619	France	Female	42	2	0.00		1	1
1	608	Spain	Female	41	1	83807.86		1	0
2	502	France	Female	42	8	159660.80		3	1
3	699	France	Female	39	1	0.00		2	0
4	850	Spain	Female	43	2	125510.82		1	1

◀ ▶

9. Categorical Data Encoding

In [17]: `df['Geography'].unique()`

Out[17]: `array(['France', 'Spain', 'Germany'], dtype=object)`

In [18]: `df = pd.get_dummies(df, drop_first = True)`

In [19]: `df.head()`

Out[19]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00		1	1	101348.88
1	608	41	1	83807.86		1	0	112542.58
2	502	42	8	159660.80		3	1	113931.51
3	699	39	1	0.00		2	0	93826.63
4	850	43	2	125510.82		1	1	79084.10

◀ ▶

10. Checking Imbalanced Data and Visualising

In [20]: `df['Exited'].value_counts()`

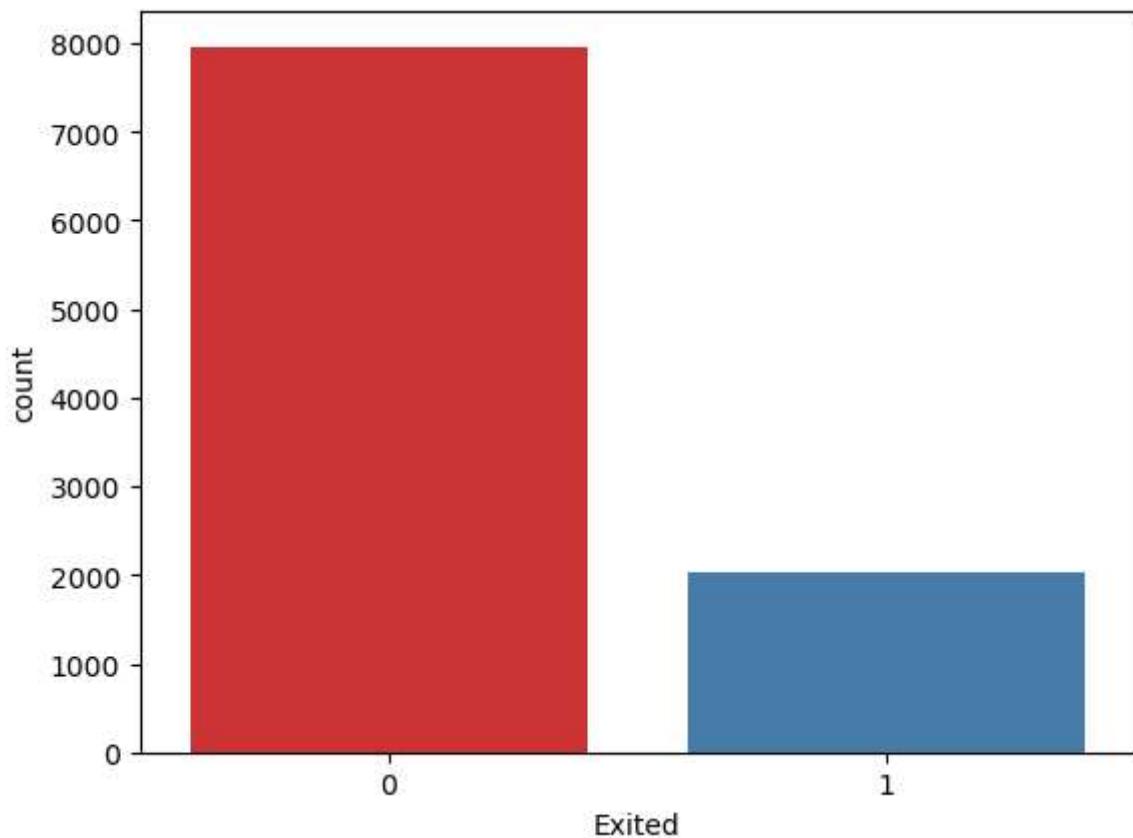
Out[20]:

Exited	count
0	7963
1	2037

Name: count, dtype: int64

In [24]: `sns.countplot(x = 'Exited', data = df, palette = 'Set1')`

Out[24]:



11. Assigning X and y Values

```
In [26]: X = df.drop('Exited', axis = 1)  
y = df['Exited']
```

```
In [27]: X
```

Out[27]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	10134
1	608	41	1	83807.86	1	0	1	11254
2	502	42	8	159660.80	3	1	0	11393
3	699	39	1	0.00	2	0	0	9382
4	850	43	2	125510.82	1	1	1	7908
...
9995	771	39	5	0.00	2	1	0	9627
9996	516	35	10	57369.61	1	1	1	10169
9997	709	36	7	0.00	1	0	1	4208
9998	772	42	3	75075.31	2	1	0	9288
9999	792	28	4	130142.79	1	1	0	3819

10000 rows × 11 columns

In [28]:

y

```
Out[28]: 0      1
         1      0
         2      1
         3      0
         4      0
         ..
        9995    0
        9996    0
        9997    1
        9998    1
        9999    0
```

Name: Exited, Length: 10000, dtype: int64

12. Balancing The Data With SMOTE

In [29]:

`from imblearn.over_sampling import SMOTE`

In [30]:

`Xs, ys = SMOTE().fit_resample(X,y)`

In [32]:

`ys.value_counts()`

```
Out[32]: Exited
          1    7963
          0    7963
```

Name: count, dtype: int64

13. Splitting The Data Into Training And Testing

```
In [33]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(Xs, ys, test_size = 0.2)
```

14. Converting Categorical Values Into Numerical Values Using The StandardScaler

```
In [34]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()
```

```
In [37]: X_train
```

```
Out[37]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Estimate
2482	567	40	4	118628.800000		1	0	0 919
12074	601	34	1	102978.844759		1	0	0 1796
9705	733	36	1	0.000000		2	0	1 1083
7505	659	27	4	0.000000		2	1	0 993
8510	550	40	8	150490.320000		1	0	0 1664
...
4813	517	39	3	0.000000		2	0	1 124
1433	755	62	1	127706.330000		2	0	1 1423
6117	850	37	2	0.000000		2	1	0 1199
13060	548	53	6	0.000000		1	0	1 735
1619	656	18	10	151762.740000		1	0	1 1270

12740 rows × 11 columns

```
In [38]: X_test
```

Out[38]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
10333	731	41	6	106110.748773	1	1	0	1948
11542	631	40	2	107841.715476	1	1	0	525
15585	429	51	4	91828.239843	1	1	0	1440
8957	585	37	6	152496.820000	1	1	1	999
334	626	37	6	108269.370000	1	1	0	55
...
12563	748	58	5	146941.760593	1	1	0	418
2450	589	46	10	107238.850000	2	1	0	370
7712	545	30	3	0.000000	2	1	0	1703
15619	624	35	1	83243.639912	2	0	1	727
5153	695	29	9	0.000000	2	1	0	1115

3186 rows × 11 columns



In [39]: y_train

```
Out[39]: 2482    0
12074    1
9705    0
7505    0
8510    1
...
4813    0
1433    0
6117    0
13060   1
1619    0
Name: Exited, Length: 12740, dtype: int64
```

In [40]: y_test

```
Out[40]: 10333   1
11542   1
15585   1
8957    0
334     0
...
12563   1
2450    0
7712    0
15619   1
5153    0
Name: Exited, Length: 3186, dtype: int64
```

```
In [41]: X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

In [42]: X_train

```
In [42]: array([[-0.88794153, -0.09917398, -0.31405339, ... , -0.75481043,
   1.58571207,  0.81529544],
 [-0.51755224, -0.70035232, -1.41921135, ... ,  1.32483596,
 -0.63063151, -1.22654924],
 [ 0.92042969, -0.49995954, -1.41921135, ... , -0.75481043,
  1.58571207,  0.81529544],
 ... ,
 [ 2.19500458, -0.39976315, -1.05082536, ... , -0.75481043,
  1.58571207,  0.81529544],
 [-1.09492378,  1.2033791 ,  0.42271858, ... , -0.75481043,
  1.58571207,  0.81529544],
 [ 0.08160689, -2.30349457,  1.89626253, ... , -0.75481043,
 -0.63063151,  0.81529544]])
```

```
In [43]: X_test
```

```
Out[43]: array([[ 8.98642084e-01,  1.02241215e-03,  4.22718582e-01, ... ,
   1.32483596e+00, -6.30631513e-01,  8.15295438e-01],
 [-1.90738169e-01, -9.91739783e-02, -1.05082536e+00, ... ,
  1.32483596e+00, -6.30631513e-01,  8.15295438e-01],
 [-2.39128628e+00,  1.00298632e+00, -3.14053391e-01, ... ,
 -7.54810432e-01,  1.58571207e+00,  8.15295438e-01],
 ... ,
 [-1.12760519e+00, -1.10113788e+00, -6.82439377e-01, ... ,
 -7.54810432e-01, -6.30631513e-01,  8.15295438e-01],
 [-2.66994787e-01, -6.00155931e-01, -1.41921135e+00, ... ,
 -7.54810432e-01,  1.58571207e+00,  8.15295438e-01],
 [ 5.06465193e-01, -1.20133427e+00,  1.52787654e+00, ... ,
 -7.54810432e-01, -6.30631513e-01, -1.22654924e+00]])
```

15. Applying Logistic Regression And Checking Accuracy, Precision Recall and F1 Scores

```
In [44]: from sklearn.linear_model import LogisticRegression
```

```
lreg = LogisticRegression()
lreg.fit(X_train, y_train)
```

```
Out[44]: ▾ LogisticRegression
LogisticRegression()
```

```
In [45]: y_pred1 = lreg.predict(X_test)
```

```
In [48]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [59]: print('Accuracy Score of SVC is: ', accuracy_score(y_test, y_pred1))
print('Precision Score of SVC is: ', precision_score(y_test, y_pred1))
print('Recall Score score of SVC is: ', recall_score(y_test, y_pred1))
print('F_1 Score of SVC is: ', f1_score(y_test, y_pred1))
```

```
Accuracy Score of SVC is: 0.77338355304457
Precision Score of SVC is: 0.7681874229346486
Recall Score score of SVC is: 0.7826633165829145
F_1 Score of SVC is: 0.775357809583074
```

In []: # With acc_score - % corresponding, pc - tp/(tp+fp), rc - tp/(tp+fn)

16. Applying SVC and Checking Score

In [60]:

```
from sklearn import svm

svm = svm.SVC()

svm.fit(X_train, y_train)
```

Out[60]:

▼ SVC
SVC()

In [61]:

```
y_pred2 = svm.predict(X_test)
```

In [62]:

```
print('Accuracy Score of SVC is: ', accuracy_score(y_test, y_pred2))
print('Precision Score of SVC is: ', precision_score(y_test, y_pred2))
print('Recall Score score of SVC is: ', recall_score(y_test, y_pred2))
print('F_1 Score of SVC is: ', f1_score(y_test, y_pred2))
```

```
Accuracy Score of SVC is: 0.8380414312617702
Precision Score of SVC is: 0.8453145057766367
Recall Score score of SVC is: 0.8272613065326633
F_1 Score of SVC is: 0.8361904761904762
```

17. Applying KNeighbors Classifier

In [64]:

```
from sklearn.neighbors import KNeighborsClassifier

KnClass = KNeighborsClassifier()
KnClass.fit(X_train, y_train)
```

Out[64]:

▼ KNeighborsClassifier
KNeighborsClassifier()

In [65]:

```
y_pred3 = KnClass.predict(X_test)
```

In [66]:

```
print('Accuracy Score of SVC is: ', accuracy_score(y_test, y_pred3))
print('Precision Score of SVC is: ', precision_score(y_test, y_pred3))
print('Recall Score score of SVC is: ', recall_score(y_test, y_pred3))
print('F_1 Score of SVC is: ', f1_score(y_test, y_pred3))
```

```
Accuracy Score of SVC is:      0.8195229127432517
Precision Score of SVC is:    0.8110091743119267
Recall Score score of SVC is: 0.8329145728643216
F_1 Score of SVC is:          0.8218159281066006
```

18. Applying Decision Tree Classifier

```
In [67]: from sklearn.tree import DecisionTreeClassifier
DTClass = DecisionTreeClassifier()
```

```
In [68]: DTClass.fit(X_train, y_train)
```

```
Out[68]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [70]: y_pred4 = DTClass.predict(X_test)
```

```
In [71]: print('Accuracy Score of SVC is:      ', accuracy_score(y_test, y_pred4))
print('Precision Score of SVC is:     ', precision_score(y_test, y_pred4))
print('Recall Score score of SVC is:  ', recall_score(y_test, y_pred4))
print('F_1 Score of SVC is:            ', f1_score(y_test, y_pred4))
```

```
Accuracy Score of SVC is:      0.8016321406151915
Precision Score of SVC is:    0.7955665024630542
Recall Score score of SVC is: 0.8115577889447236
F_1 Score of SVC is:          0.8034825870646767
```

19. Applying Random Forest

```
In [72]: from sklearn.ensemble import RandomForestClassifier
RFClass = RandomForestClassifier()
```

```
In [73]: RFClass.fit(X_train, y_train)
```

```
Out[73]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [74]: y_pred5 = RFClass.predict(X_test)
```

```
In [75]: print('Accuracy Score of SVC is:      ', accuracy_score(y_test, y_pred5))
print('Precision Score of SVC is:     ', precision_score(y_test, y_pred5))
print('Recall Score score of SVC is:  ', recall_score(y_test, y_pred5))
print('F_1 Score of SVC is:            ', f1_score(y_test, y_pred5))
```

```
Accuracy Score of SVC is:      0.8706842435655995
Precision Score of SVC is:    0.8701380175658721
Recall Score score of SVC is: 0.8712311557788944
F_1 Score of SVC is:          0.8706842435655996
```

20. Applying Gradient Boost Classifier

```
In [76]: from sklearn.ensemble import GradientBoostingClassifier
        GBClass = GradientBoostingClassifier()
```

```
In [77]: GBClass.fit(X_train, y_train)
```

```
Out[77]: GradientBoostingClassifier()
          GradientBoostingClassifier()
```

```
In [78]: y_pred6 = GBClass.predict(X_test)
```

```
In [79]: print('Accuracy Score of SVC is:', accuracy_score(y_test, y_pred6))
        print('Precision Score of SVC is:', precision_score(y_test, y_pred6))
        print('Recall Score score of SVC is:', recall_score(y_test, y_pred6))
        print('F_1 Score of SVC is:', f1_score(y_test, y_pred6))
```

```
Accuracy Score of SVC is: 0.8402385436283741
Precision Score of SVC is: 0.8429385687143762
Recall Score score of SVC is: 0.8360552763819096
F_1 Score of SVC is: 0.8394828129927467
```

21. Saving And Visualising Accuracy And Precision Scores

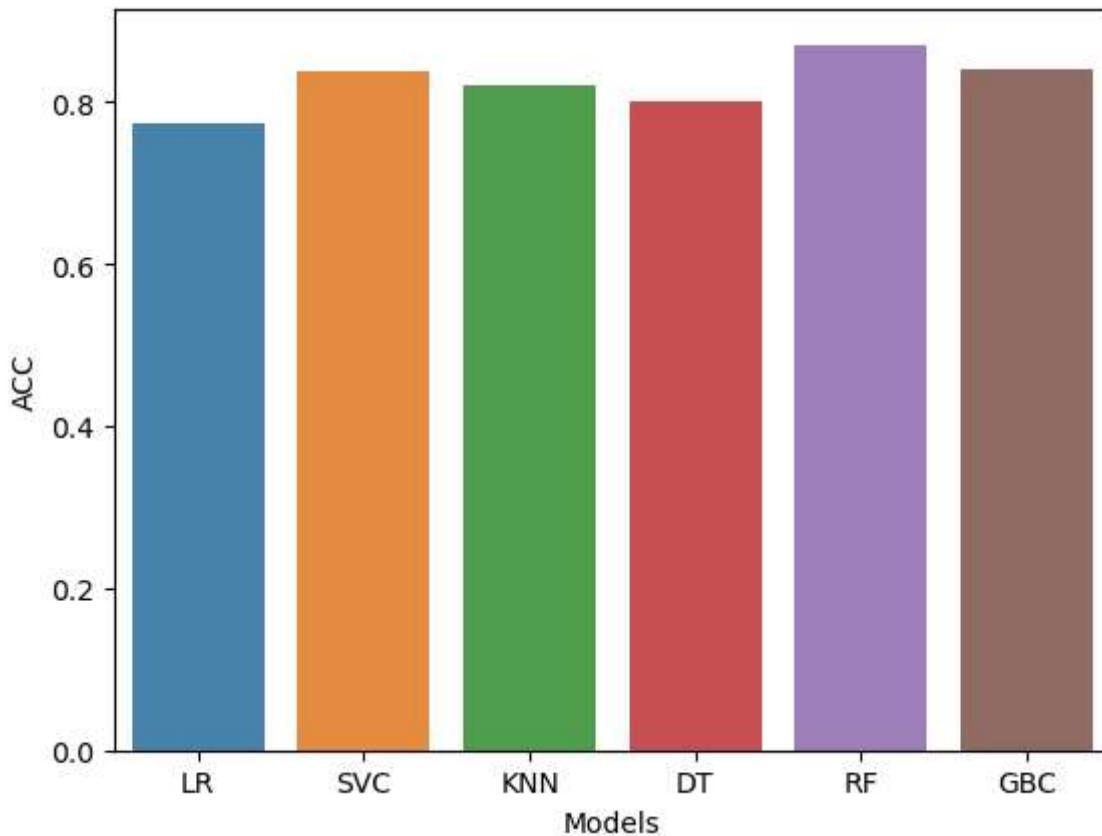
```
In [118...]: accuracies = pd.DataFrame({'Models': ['LR', 'SVC', 'KNN', 'DT', 'RF', 'GBC'], 'ACC': [0.773384, 0.838041, 0.819523, 0.801632, 0.870684, 0.840239]})
```

```
In [115...]: print(accuracies)

sns.barplot(x=accuracies['Models'], y=accuracies['ACC'], alpha=0.9 )
```

	Models	ACC
0	LR	0.773384
1	SVC	0.838041
2	KNN	0.819523
3	DT	0.801632
4	RF	0.870684
5	GBC	0.840239

```
Out[115]: <Axes: xlabel='Models', ylabel='ACC'>
```



In [119...]

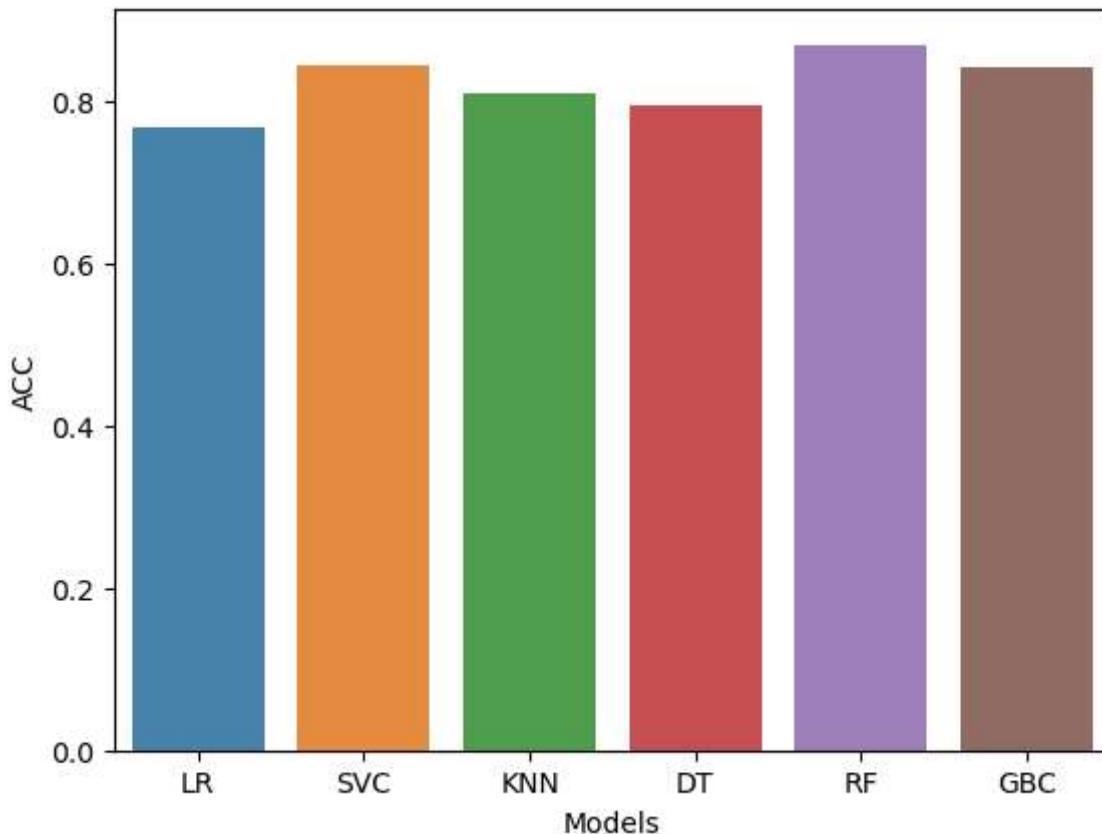
```
precisions = pd.DataFrame({'Models':['LR', 'SVC', 'KNN', 'DT', 'RF', 'GBC'], 'ACC':[0.768187, 0.845315, 0.811009, 0.795567, 0.870138, 0.842939]})  
sns.barplot(x=precisions['Models'], y=precisions['ACC'], alpha=0.9 )
```

In [120...]

```
print(precisions)  
  
sns.barplot(x=precisions['Models'], y=precisions['ACC'], alpha=0.9 )
```

	Models	ACC
0	LR	0.768187
1	SVC	0.845315
2	KNN	0.811009
3	DT	0.795567
4	RF	0.870138
5	GBC	0.842939

Out[120]:



22. Saving The Model For Later Use

```
In [123...]: Xs = sc.fit_transform(Xs)  
RFClass.fit(Xs, ys)
```

```
Out[123]: RandomForestClassifier  
RandomForestClassifier()
```

```
In [125...]: import joblib  
joblib.dump(RFClass, 'Customer_Churn_Prediction_Model')
```

```
Out[125]: ['Customer_Churn_Prediction_Model']
```

23. Loading The Model For Prediction

```
In [126...]: model = joblib.load('Customer_Churn_Prediction_Model')
```

```
In [142...]: model.predict([[608, 41, 1, 100807, 1, 0, 1, 200000.58, 0, 1, 0]])
```

```
Out[142]: array([0], dtype=int64)
```

```
In [130...]: X
```

Out[130]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	10134
1	608	41	1	83807.86	1	0	1	11254
2	502	42	8	159660.80	3	1	0	11393
3	699	39	1	0.00	2	0	0	9382
4	850	43	2	125510.82	1	1	1	7908
...
9995	771	39	5	0.00	2	1	0	9627
9996	516	35	10	57369.61	1	1	1	10169
9997	709	36	7	0.00	1	0	1	4208
9998	772	42	3	75075.31	2	1	0	9288
9999	792	28	4	130142.79	1	1	0	3819

10000 rows × 11 columns



In [143...]

ys

Out[143]:

```

0      1
1      0
2      1
3      0
4      0
..
15921  1
15922  1
15923  1
15924  1
15925  1
Name: Exited, Length: 15926, dtype: int64

```

In []:

